

5 Claude Skills

Five real, open-source skills that extend Claude past chat and coding — verified repos, exact install commands, honest costs, and what each one actually does.

5

independent, open-source projects — all verified

2

of the five need a paid third-party account or key

2

needed real identification — Section 01 explains

Sanchit Pandey

AI systems, content & automation
Lucknow, India

Five real projects, verified individually

Every skill in this document was checked against its own source repository rather than taken on faith from a reel or a listicle. Three matched their names exactly. Two needed identifying — the closest, best-documented real project matching the description, since no skill exists under that literal name. Section 01 explains which is which before you install anything.

All five are free, open-source, and independently built — none are official Anthropic products, and two require a separate paid account with a third party to actually use what they produce. That distinction matters enough to flag clearly, so it's in the table below before anything else.

WHAT THIS DOCUMENT IS

Verified install commands, what each skill genuinely does, independently-checked performance claims where they exist, and the real cost of using each one beyond the install itself.

The five, at a glance

Skill	Verified as	Cost beyond install
Ponytail	Exact name match — DietrichGebert/ponytail	Free
Claude Mem	Exact name match — thedotmack/claude-mem	Free
Obsidian Skill	Closest match — eugeniughelbur/obsidian-second-brain	Free core; optional paid API keys for 7 research commands
Banana Pro Director	Closest match — AgriciDaniel/banana-claude	Free skill; needs a Google AI Studio API key (free tier available)
Cinema World Builder	Closest match — the cinema-worldbuilder skill	Free skill; produces prompts for Higgsfield, a paid third-party video platform

A skill is software, not a setting

Worth repeating for any list like this: a Claude Code skill can read your files and, depending on what it's built to do, run commands or call external APIs. Treat installing one the way you'd treat installing any other software written by someone you don't know.

- ☐ Open the repository before installing and skim the README — all five below are linked with exact repo names for this reason
- ☐ Prefer the plugin marketplace install shown for each skill over any script that pipes straight into bash
- ☐ Check who maintains it and how active the repo is — star counts and recent commits are a signal, not a guarantee
- ☐ Note which skills ask for an API key, and only paste keys into files you've read
- ☐ Two of these five route your prompts to a third-party API (Google, or a video platform) — read that provider's own data terms if that matters to you

Ponytail

"Makes your AI agent think like the laziest senior dev in the room. The best code is the code you never wrote." — the project's own tagline.

What it actually does

Before writing code, Ponytail forces the agent through a seven-rung ladder — does this need to exist at all, is it already in the codebase, does the standard library do it, is there a native platform feature, is there an installed dependency, can it be one line — and only then writes the minimum that works. It never touches validation, error handling, security, or accessibility.

The numbers, honestly

Source	Result
Project's own agentic benchmark (real Claude Code session, real FastAPI+React repo, 12 tickets)	-54% lines of code, -22% tokens, -20% cost, -27% time, 100% safety
Independent re-benchmark by JetBrains (own test set, 80 paired tasks)	-15% code, -10.3% cost (p=0.004) — more conservative, on different tasks
Earlier single-shot benchmark (superseded)	80-94% less code — since flagged by the project itself as an inflated baseline artifact and corrected

The gap between the two isn't dishonesty — Ponytail's own maintainer publicly corrected the inflated early number, and JetBrains tested on different, more conservative tasks. Expect real but moderate savings, biggest on tasks where an agent tends to over-build.

Install (Claude Code)

TERMINAL

```
/plugin marketplace add DietrichGebert/ponytail
/plugin install ponytail@ponytail

# Two separate messages - the install needs both sent one after
# the other, not combined.
```

Commands: `/ponytail [lite|full|ultra|off]`, `/ponytail-review`, `/ponytail-audit`, `/ponytail-gain` (shows your own measured impact).

Claude Mem

Persistent memory for Claude Code: captures what happens in a session, compresses it with AI, and injects the relevant parts back into future sessions.

What it actually does

By default, every new Claude Code session starts from zero — project conventions, past decisions, and prior context are gone unless they're written into `CLAUDE.md` by hand. Claude Mem automates that: it watches tool usage during a session, generates a semantic summary, stores it in local SQLite, and surfaces the relevant pieces the next time they're needed — without you re-explaining anything.

What makes it more than a transcript dump

- Progressive disclosure — retrieves layered context with visible token cost, rather than dumping everything back in and burning your context window
- Semantic search — query past sessions by meaning (`mem:search`), not just exact keyword match
- Privacy control — wrap anything in `<private>` tags during a session and it's excluded from storage entirely
- A web viewer for watching the memory stream in real time, plus a Claude Desktop skill so you can search project memory from outside the terminal too

Install (Claude Code)

TERMINAL

```
/plugin marketplace add thedotmack/claude-mem  
/plugin install claude-mem@thedotmack
```

Commands: `mem:search`, `mem:recall`, `mem:forget`, `mem:stats`. Apache-2.0 licensed, storage is local SQLite — nothing leaves your machine by default.

The Obsidian Skill

No single skill is officially called "the Obsidian Skill" — several good ones exist. The best-documented and most actively maintained as of this writing is `obsidian-second-brain` by Eugeniu Ghelbur.

Extends Andrej Karpathy's "LLM Wiki" pattern: instead of only appending new notes, it rewrites existing pages as new information arrives, and resolves contradictions automatically.

What it actually does — 44 commands across four layers

Layer	Does	Example command
Operations (28)	Saves, ingests, organises and maintains the vault	<code>/obsidian-save</code> , <code>/obsidian-ingest</code>
Thinking (8)	Challenges your ideas and surfaces patterns you didn't name	<code>/obsidian-challenge</code> , <code>/obsidian-emerge</code>
Context (1)	Loads your identity and current state at session start	<code>/obsidian-world</code>
Research (7)	Pulls live external knowledge into the vault	<code>/research</code> , <code>/x-pulse</code> , <code>/youtube</code>

Two things worth knowing before installing

- The 37 core commands need no API key at all — only the 7 research commands do, and even those fall back to free sources (Wikipedia, arXiv, Reddit and others) when no key is set
- Your vault stays plain markdown — the project's own pitch is that this is the anti-lock-in alternative to closed SaaS tools like Notion AI or Mem: stop using the skill any time and you still have your full vault

Install (Claude Code)

TERMINAL

```
/plugin marketplace add eugeniughelbur/obsidian-second-brain
/plugin install obsidian-second-brain@obsidian-second-brain

# Then add your vault path to ~/.claude/settings.json:
# "env": { "OBSIDIAN_VAULT_PATH": "/path/to/your/vault" }
# Restart Claude Code, then run /obsidian-init inside the vault.
```

MIT licensed. If you'd rather have simple direct read/write into your vault without 44 commands, the lighter alternative is an Obsidian MCP server paired with the free community "Local REST API" Obsidian plugin — worth knowing as an option if this full system is more than you need.

Banana Pro Director

No skill is named exactly this. The closest real match — same concept, same "Director" framing — is banana-claude by AgriciDaniel: "Claude acts as Creative Director" over Google's Nano Banana (Gemini) image models.

Unlike a simple API wrapper, Claude interprets your intent, applies Google's official five-component prompt formula, and iterates — generating, looking at the result, and re-generating until it holds up.

What it actually does

- Turns a plain request ("a hero image for a coffee shop website") into a fully art-directed prompt before it ever reaches the image model
- Generates, edits, and batch-produces variations — up to 4096×4096, 14 aspect ratios including cinematic 21:9
- Maintains character and style consistency across a multi-turn creative session
- Ships a browsable database of 2,500+ reference prompts (`/banana inspire`) for when you want a starting point, not a blank page

Install (Claude Code)

TERMINAL

```
/plugin marketplace add AgriciDaniel/banana-claude
/plugin install banana-claude@banana-claude-marketplace

# Needs a free Google AI Studio API key: aistudio.google.com
```

Commands: `/banana generate "...", /banana edit photo.png "...", /banana chat, /banana inspire`.

The skill itself is free; the Gemini image API it calls is not unlimited. Google AI Studio's free tier covers light use — generating dozens of images in one session will likely draw on paid usage. Check current Gemini API pricing before a large batch run.

Cinema World Builder

Matches an existing real skill closely: cinema-worldbuilder, one of four skills in an "ai-film-pipeline" set (image prompts, Seedance video prompts, orchestrator, QC) for high-end AI video production on the Higgsfield platform.

"Cinema worldbuilding director for Seedance video prompts with locked compositional rigor and mechanical prompt precision." — the skill's own description.

What it actually does

This skill does not generate video itself — it generates the detailed, structured prompt that Seedance (the video model inside Higgsfield) needs to produce a cinematic shot. It uses a five-mode grammar (Narrative, Studio, Action, Performance, Atmospheric) and assembles a fixed sequence of blocks every time: Scene & Mood, Frame Map, Subject Lock, Cross-Frame Rules, Movement, Last Frame, World Plate, Sound Bed, Capture Realism, Camera Capture — with your reference images tagged and anchored to exact points in the prompt.

The honest cost picture

- The skill itself is free and installs like any other
- Higgsfield is not. It's a separate, paid AI video platform — a real company that raised an \$80M extension in January 2026 at a \$1.3B valuation. Running the prompts this skill writes requires an active Higgsfield account
- This is a prompt engineer, not a renderer. Its output is a copy-paste block for Higgsfield's Seedance interface, not a finished video file

How it fits with Banana Pro Director

The pairing in the reel is accurate: generate reference images with the previous skill, then feed those images into this one to produce the cinematic video prompt that animates them — image generation and video-prompt engineering as two separate, deliberately scoped steps.

Multiple public forks of this skill set exist under slightly different repository owners with identical content — check whichever copy you install for a recent commit date before relying on it.

Using more than one together

These five weren't built as a set, but they don't conflict — they sit at different layers.

Skill	Layer it operates at
Ponytail	How code gets written, on every coding task
Claude Mem	What Claude remembers between sessions, on every task
Obsidian Skill	Your personal knowledge base — separate from coding work entirely
Banana Pro Director	Image generation, invoked when you need a visual
Cinema World Builder	Video-prompt generation, invoked after you have reference images

A realistic combined session

Ponytail and Claude Mem can both run silently in the background of an ordinary coding session. Separately, for a content project: Banana Pro Director generates a set of reference images, Cinema World Builder turns them into a Higgsfield-ready video prompt, and Obsidian Second Brain logs the whole creative decision trail to your vault via `/obsidian-save` so you're not starting from zero on the next similar project.

Every installed skill's description sits in context on every turn, whether triggered or not. Five well-scoped skills is a reasonable working set. Install what you'll actually use — not all five because the list existed.

Sources & notes

Checked August 2026, directly against each project's own repository. Skill ecosystems move fast — star counts, exact commands, and pricing all drift; re-check the repo before relying on specifics months from now.

What the guide says	Source
Ponytail mechanism, agentic benchmark numbers, and the corrected single-shot figure	github.com/DietrichGebert/ponytail — official repo and benchmark results
JetBrains' independent re-benchmark of Ponytail	blog.jetbrains.com , July 2026
Claude Mem mechanism, commands, licensing	clausedirectory.org and alternativeto.net listings, cross-checked against thedotmack.com/claude-mem
obsidian-second-brain's 44 commands, install steps, no-key core, Karpathy LLM Wiki lineage	github.com/eugeniughelbur/obsidian-second-brain — official repo, README
banana-claude's Creative Director framing, commands, install, API key requirement	github.com/AgriciDaniel/banana-claude — official repo
cinema-worldbuilder's mechanism, five-mode grammar, block structure	SKILL.md from the ai-film-pipeline repository, cross-checked across two public forks
Higgsfield's funding, valuation, and status as a separate paid platform	Wikipedia, citing Reuters and TechCrunch reporting, January 2026

None of the five skills profiled here are official Anthropic products. This is an independent guide, not sponsored by or affiliated with any of the projects or companies named.

THAT'S THE FIVE

Verify the name. Then verify the cost.

Most "5 skills that change everything" lists ask you to trust the names. Two of these five needed real identification work before they could be installed correctly, and two come with a cost the original pitch never mentions. None of that makes them less worth using — it makes them worth using deliberately.

Start with Ponytail and Claude Mem — free, no external accounts, useful on almost every coding session. Add the other three once you know which creative workflow you're actually building.

Sanchit Pandey

AI systems, automation and content.

If you find the actual skill this list was pointing at and it isn't the one profiled here, send it over — a corrected identification is worth more than a confident guess.

© 2026 Sanchit Pandey. Share it freely — keep the attribution.