
A PRACTICAL GUIDE · STRAWBERRY BROWSER · 2026

Getting Started With Strawberry

What the self-driving browser actually does, how credits and pricing work, ready-to-use companion instructions for leads, hiring and invoices — and the safety trade-offs worth understanding before you hand it real access.

2,000

free credits every month,
no card required

6

core use cases the product
is actually built around

1

safety habit that matters more
than any setting — Section 08

Sanchit Pandey

AI systems & automation
Lucknow, India

What Strawberry actually is

Strawberry is a real, downloadable web browser — not a plugin, not a chat window bolted onto Chrome — built by the Stockholm-based company Dendrite Systems. What makes it different from Chrome or Edge is that it ships with "companions": AI agents that can see your tabs and act inside them the same way you do, by scrolling, clicking, typing and reading.

The pitch in the reel — leads, hiring, invoices — maps directly to how the product describes itself: sales prospecting, candidate sourcing, and inbox or document work are three of its six named use cases. This document explains how to actually get value from it, plus the one topic worth understanding before you connect it to anything sensitive.

THE HONEST SUMMARY

A genuinely capable browser-automation tool, free to start, with real safety features built in — and a category-wide security risk (Section 08) that applies to every product like it, not just this one. Worth using deliberately, not blindly.

The basics

Made by	Dendrite Systems (Stockholm), founded by Charles Maddock, Arian Hanifi and Sebastian Thunman
Platforms	macOS and Windows. No Linux or mobile app at the time of writing
Funding	\$6M seed, led by General Catalyst and EQT Ventures, with participation from the founders of Lovable, Hugging Face and Supabase
Cost to start	Free — 2,000 credits per month, no card required
Get it	strawberrybrowser.com

How it actually works

Three ideas explain almost everything about how you will use it day to day.

1. It's a browser first

Ordinary browsing — reading pages, watching videos, filling in a form yourself — uses no credits at all and works exactly like Chrome. You can import bookmarks, passwords and Chrome extensions during setup. Credits are spent only when you talk to a companion or a companion takes action on your behalf.

2. Companions share your tabs, not a sandbox

This is the structural difference from a chatbot. Because a companion operates inside your actual browser session, it can use whatever you are already logged into — your CRM, your inbox, LinkedIn, internal dashboards — without you re-authenticating it separately. That is also precisely why Section 08 matters: the companion's reach is your reach.

3. Two ways to use a companion

Mode	How it works	Good for
On-demand	You give an instruction now and watch it work in real time — you can pause, correct, or take over at any point	One-off research, a single outreach batch, anything you want to review closely the first few times
Routines	A trigger (a new email, a Slack mention, a PR opened) automatically starts a companion, which works and notifies you before anything irreversible happens	Recurring work you have already tested manually and trust

Start everything as on-demand. Only convert a task into a routine once you have watched it succeed several times and understand exactly what it does at each step.

Getting started

- 01** Download the desktop app from strawberrybrowser.com for macOS or Windows.
- 02** Import your data if you want to — Chrome extensions, bookmarks and passwords can be carried over. This is optional; you can also start clean.
- 03** Meet your companion. The onboarding introduces a default personal companion and a set of pre-built ones for common roles.
- 04** Try one pre-built use case before writing a custom companion — it is the fastest way to learn what "good instructions" look like in this product specifically.
- 05** Check your credit balance as you go, so pricing (Section 03) does not surprise you mid-task.

What you can hand it, in the product's own words

Research companies, source leads on LinkedIn, update CRMs, transcribe meetings, create slide decks, pull data into spreadsheets, draft emails, and run recurring browser work — anything that can be done in a browser.

The six named use cases

Use case	What it does
Market Research	Profiles companies and sizes up markets using Crunchbase, Statista, G2 and LinkedIn
Sales Prospecting	Finds LinkedIn leads, enriches CRM records, drafts personalised outreach
Founder / Ops Updates	Pulls metrics from Sheets, Slack and docs into a written update
Candidate Sourcing	Searches LinkedIn and job boards, messages qualified candidates
Social Listening	Tracks what people say about you across Reddit, X and review sites
Web Data Extraction	Pulls pricing, listings or company data into a structured sheet

It also handles meeting transcription without a bot joining the call, and can build slide decks, posters and dashboards from sources it finds on the web, including behind login walls.

Pricing and credits

Understanding credits before you start saves you from an unpleasant surprise on a big task.

Plan	Price	Credits / month
Free	₹0	2,000
Intern	\$20/mo	8,000
Part-time	\$100/mo	30,000
Full-time	\$250/mo	75,000
Top-up (any plan)	\$10 per 1,000 credits	As needed

What actually spends a credit

- Regular browsing spends nothing — reading, watching, manually filling a form
- Chatting with a companion spends credits, roughly in proportion to how much it reads and reasons through
- A companion taking action — clicking, typing, navigating across pages — spends credits for each step, so a 20-site research task costs meaningfully more than a 2-site one

Budgeting a task before you run it

- 01** Run a small version first — 3 leads instead of 30, one candidate search instead of ten — and note the credit cost.
- 02** Multiply up for the real batch size, and compare against your plan's monthly allowance.
- 03** For anything recurring, check the cost of one run before turning it into an automated routine that fires repeatedly without you watching.

The free tier is genuinely enough to evaluate whether this fits your workflow before spending anything — use it to test the three companion types in Section 05 before committing to a paid plan.

Writing instructions a companion can actually follow

The gap between a companion that works well and one that gets stuck is almost always in how specific the instructions are — the same lesson as prompting any AI tool, with one addition: browser actions need explicit boundaries.

Five rules

- 01** Name the exact source. "Find leads" is vague. "Find leads on LinkedIn Sales Navigator, filtered to VP of Sales at Series B companies" is executable.
- 02** Give a hard stopping point. A number, a time limit, or a condition — "find 20", "spend no more than 10 minutes", "stop after checking 3 sources". Without one, a companion can wander.
- 03** State what counts as done. "Add each lead to the sheet with name, title, company and LinkedIn URL" is checkable. "Research some leads" is not.
- 04** Say what NOT to do, explicitly, for anything irreversible — "do not send anything, only draft" or "do not submit the form, just fill it and stop".
- 05** Ask for a summary at the end, not just the completed action — "tell me what you did and flag anything you were unsure about" turns a black box into something you can actually review.

The instruction quality matters more here than with a chat-only AI tool, because a vague instruction does not just produce a mediocre answer — it can produce the wrong action, taken for real, in a real account.

Three companions to start with

The exact three use cases from the reel — leads, hiring, invoices — written as instructions you can copy and adapt.

Companion 1 — Lead finder

COMPANION INSTRUCTIONS

Find 15 leads matching this profile: [JOB TITLE] at [COMPANY TYPE / INDUSTRY], company size [RANGE], location [WHERE].

Search LinkedIn. For each lead, add a row to [SHEET NAME] with: name, title, company, company size, LinkedIn URL, and one line on why they fit this profile.

Do not send any messages or connection requests. Only research and log. Stop once you have 15 that genuinely match - do not pad the list with weak fits to hit the number.

When done, summarise: how many you found, any pattern in who did NOT fit, and whether the search criteria should be adjusted.

Companion 2 — Candidate shortlist

COMPANION INSTRUCTIONS

Search LinkedIn and [JOB BOARD] for candidates matching: [ROLE], [YEARS OF EXPERIENCE], [KEY SKILLS], open to new roles, based in [LOCATION or remote].

Shortlist up to 10. For each, log: name, current role and company, years of relevant experience, key skills matched, profile URL, and a one-line note on fit.

Flag anyone who looks like a strong fit but is missing one of the criteria, rather than silently excluding them.

Do not message anyone. Draft an outreach note for my review separately once I've seen the shortlist.

Companion 3 — Invoice and inbox triage

COMPANION INSTRUCTIONS

Scan my inbox for invoices and payment reminders from the last 30 days. For each one found, log: sender, amount, due date, invoice number, and whether it appears already paid (check for a reply thread confirming payment).

Flag anything overdue, or due within 5 days, at the top of the list.

Do not reply to any email or mark anything as read/unread. Only read and log. Tell me if any email looks like it could be a fake invoice or phishing attempt rather than a real one.

Note the shared pattern: every one of these explicitly forbids the irreversible action (sending, messaging, replying) and asks for a summary at the end. Keep that pattern for anything you write yourself.

Turning a task into a routine

Once a companion's on-demand version works reliably, Strawberry can trigger it automatically — "reply to leads and log them to the CRM" when a labelled email arrives, or "draft a reply" when someone mentions you in a Slack channel.

Before automating anything

- ☐ Run the exact task on-demand at least 5 times with no correction needed
- ☐ The instruction has an explicit stopping point and forbids the irreversible step
- ☐ You have reviewed what happens when it encounters something unexpected — does it ask, or guess?
- ☐ The trigger is specific enough that it will not fire on the wrong emails or messages
- ☐ You know how to pause or disable the routine quickly if something looks wrong

Example routine structures

Trigger	Action	Where it stops
Gmail: label "Leads"	Draft a reply and log the lead to the CRM	Notifies you before sending; you approve
Slack: mention in #design	Reads the thread, drafts a reply	Posts only after you confirm
GitHub: PR opened	Reproduces the failing test, opens a fix PR	Opens the PR; does not merge

A routine is a companion you are no longer watching. Everything about its instructions should be stricter than the on-demand version, not looser.

What Strawberry does for safety

Worth knowing what is actually built in, separate from the habits in Section 08 that are on you.

Feature	What it means
You're in the driver's seat	Companions ask for confirmation before taking irreversible actions — sending, paying, deleting, submitting
SOC 2 & GDPR	The company states it meets SOC 2 requirements for secure data handling
Passwords stay local	Encrypted on your device; the company states they are never visible to Strawberry or to the underlying AI providers
No training on your data	The company states your browsing activity is not used to train AI models

What this does and does not cover

"Asks before irreversible actions" is a real, useful guardrail — and it is also exactly the guardrail that the security research in the next section is about, because the whole point of a prompt injection attack is tricking the companion into treating a malicious instruction as though it came from you, sidestepping your intent rather than the confirmation dialog itself. The built-in safety features and the habits in Section 08 are not redundant — you need both.

The risk worth understanding: prompt injection

This is not a Strawberry-specific flaw. It is a structural, industry-wide risk in every agentic browser as of 2026 — including the best-known ones from the largest AI companies. Understanding it is what lets you use any of these tools well.

What it is

An agentic browser reads two things at once when it acts for you: your instruction, and the content of the page it is looking at. Prompt injection is when an attacker hides a second set of instructions inside that page — in white-on-white text, a spoiler tag, a hidden comment — written to look like a command. The browser cannot always reliably tell your real instruction apart from text on the page pretending to be one.

Why it is taken seriously in 2026

- Security researchers demonstrated working prompt injection attacks across every major agentic browser tested at a Black Hat USA 2026 session, including products from large, well-resourced companies
- A University of Washington study of seven popular agentic browsers found four allowed attackers to bypass the same-origin policy — the fundamental rule that stops one website from reading another's data — via prompt injection
- OpenAI's own security team has stated that prompt injection is one of the most significant risks they actively defend against, and continues shipping hardening updates against new attack classes
- The UK's National Cyber Security Centre and OpenAI have both said the risk likely cannot be fully "solved" with current techniques — it is structural to how these agents read the web, not a bug in one product
- Gartner has advised some enterprises to restrict or block agentic browsers for sensitive work while the category matures

What this means practically

It does not mean avoid the category. It means treat a browser companion the way you would treat a new employee with your logins: genuinely useful, worth delegating to, and not handed your most sensitive access on day one without a reason to trust it there specifically.

Practical habits, not paranoia

None of these make the tool less useful. They make the risk proportionate to what you are actually doing.

- ☐ Keep banking, payment and highly sensitive accounts logged out in the browser a companion uses, or in a separate browser profile it cannot touch
- ☐ Require confirmation for anything that sends, pays, deletes or submits — even if the product would let you turn that off, don't
- ☐ Be more cautious with companions that browse open, untrusted pages (research, social listening) than ones working only inside your own tools (your CRM, your inbox)
- ☐ Review a routine's activity log periodically, not just when something looks wrong
- ☐ If a companion's behaviour suddenly does not match what you asked, stop it and check what page it was reading when it happened
- ☐ Keep the browser and the product itself updated — security hardening for this category ships frequently and matters
- ☐ Use a separate, less-privileged login for a companion's daily work where the product supports it, rather than your primary admin account

This category will keep getting safer — every major vendor is actively working the problem. Being informed now costs you five minutes and means you are using the tool deliberately instead of by default.

How it compares

A fair, brief placement — not a ranking. The right tool depends on what you are already using.

	What it is	Where it's distinct
Strawberry	A standalone browser with role-based "companions" and credit-based pricing	Purpose-built for delegated multi-step work — sales, sourcing, research — rather than general browsing assistance
Claude in Chrome	A browsing agent inside Chrome, from Anthropic	Runs inside your existing Chrome setup rather than replacing the browser
ChatGPT Atlas	OpenAI's own agentic browser	Deep ChatGPT integration; actively publishes its own prompt-injection hardening research
Perplexity Comet	Search-company-built agentic browser	Strong at research and summarisation-led tasks

None of these fully solve the risk in Section 08 — it is a property of the category, not a gap in any one product. Choose based on which tools you already use daily and want the agent working inside, and apply the same habits from Section 09 regardless of which you pick.

Sources & notes

Product details checked directly against strawberrybrowser.com, August 2026. Pricing, credit allowances and platform support change — verify on the product's own pricing page before making a purchasing decision.

What the guide says	Source
Product description, use cases, pricing, credits, safety features, FAQ	strawberrybrowser.com, official product site
Founders (Charles Maddock, Arian Hanifi, Sebastian Thunman), \$6M seed round, investors	Silicon Canals and JustAINews funding coverage, October 2025
Public open beta launch	Hyperight, June 2026
Every agentic browser tested at Black Hat USA 2026 showed prompt injection vulnerabilities	Dark Reading, coverage of Brave Software's "Attacking and Defending AI Browsers" session
University of Washington study: 4 of 7 agentic browsers allow same-origin-policy bypass via prompt injection	UW News, June 2026
OpenAI's ongoing hardening work against prompt injection in Atlas	OpenAI, "Continuously hardening ChatGPT Atlas against prompt injection attacks"
UK NCSC and OpenAI position that prompt injection cannot currently be fully solved	Multiple 2026 security-press summaries citing both organisations

This document is an independent guide, not an official Strawberry or Dendrite Systems resource, and not sponsored by them. Nothing here is security advice for a specific high-stakes deployment — for that, consult a security professional.

THAT'S THE GUIDE

Delegate the work. Keep the judgement.

A browser that can act for you is a genuine step change for repetitive work — leads, sourcing, research, inbox triage all get meaningfully faster. It is also, structurally, a wider door into your accounts than a browser has ever been before. Both of those things are true at once, and using it well just means acting like it.

Start on the free plan, with the three companions in Section 05, on tasks that only read and log — nothing that sends or pays. Earn your way up to routines once you have watched it work.

Sanchit Pandey

AI systems and automation.

If you build a companion worth sharing, or hit a prompt-injection-shaped surprise, tell me — both are exactly what future versions of this guide need.

© 2026 Sanchit Pandey. Share it freely — keep the attribution. Independent guide, not affiliated with Dendrite Systems.